

# Assignment-2

Shivansh Garg (2024MCS2450) Rohit Patidar (2024JCS2042)

April 2025

## Introduction

This document presents an in-depth explanation of the control flow mechanisms and scheduling strategy implemented in a custom shell environment. It focuses on how different key combinations (`Ctrl+C`, `Ctrl+B`, `Ctrl+F`, and `Ctrl+G`) trigger specific signals and the way these signals are handled within the shell and kernel space.

The first part provides a detailed step-by-step breakdown of the signal handling mechanism, illustrating how each key combination affects process states and scheduling behavior. This includes the termination signal (`SIGINT`), background suspension (`SIGBG`), foreground resumption (`SIGFG`), and custom user-defined signal (`SIGCUSTOM`).

The second part of the document explains the dynamic priority scheduling algorithm. The scheduling policy uses a weighted function of CPU usage and wait time, influenced by two tunable parameters:  $\alpha$  and  $\beta$ . These parameters control how a process's priority evolves over time, enabling fine-grained control over system responsiveness and fairness.

Together, these mechanisms form a responsive and balanced process management system, suitable for managing multiple types of workloads in a multi-process environment.

## Individual Control Flow Explanation

This section explains the control flow for each signal triggered by the key combinations `Ctrl+C`, `Ctrl+B`, `Ctrl+F`, and `Ctrl+G`.

### 1. `Ctrl+C` $\rightarrow$ `SIGINT` (Terminate)

1. **User Input:** User presses `Ctrl+C`.

#### 2. Shell Input Handling:

- The shell detects `Ctrl+C` and interprets it as a command to terminate processes.
- It calls `send_signal_to_all(SIGINT);` to broadcast the termination signal.

#### 3. Inside `send_signal_to_all()`:

- Iterates through all processes.
- For each process (except shell and init):
  - Sets `p->killed = 1`
  - If state is `SLEEPING`, `SUSPENDED`, etc., sets `p->state = RUNNABLE` so that the process wakes up and gets killed in the next scheduler cycle.

4. **Result:** Processes are terminated.

```
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ test
Hello 0
Hello 1
Hello 2
Hello 3
Hello 4
Hello 5
Ctrl-C is detected by xv6
PID: 3
TAT: 569
WT: 0
RT: 0
#CS: 0
$ █
```

Figure 1: Ctrl + C

## 2. Ctrl+B → SIGBG (Send to Background)

1. **User Input:** User presses Ctrl+B.

2. **Shell Input Handling:**

- Shell calls `send_signal_to_all(SIGBG)` ; to suspend all other processes.

3. **Inside `send_signal_to_all()`:**

- Iterates over processes (except shell and init).
- For eligible processes in RUNNING, RUNNABLE, etc.:
  - Sets `p->state = SUSPENDED`
  - Sets `p->suspended = 1`
  - Reassigns their parent to `initproc` (if `pid > 2`)

4. **Result:** All active processes are suspended and sent to background.

```

2
xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ test2
Parent started with pid: 3
Hello , I am parent
Child started with pid: 4
Hi there , I am child
Hi there , I am child
Hello , I am parent
Hi there , I am child
Hello , I am parent
Hi there , I am chilHello , I am parent
d
Ctrl-B is detected by xv6
$ Ctrl-F is detected by xv6
Hi there , I am child
Hello , I am parent
Hi there , I am child
Hello , I am parent
Hello , I am parent
Hello , I am parent
Hi there , I am child
Hello , I am parent
Hi there , I am child
Ctrl-C is detected by xv6
PID: 3
TAT: 446
WT: 292
RT: 1
#CS: 153
PID: 4
TAT: 446
WT: 293
RT: 1
#CS: 152

```

Figure 2: Ctrl + B

### 3. Ctrl+F → SIGFG (Bring to Foreground)

1. **User Input:** User presses Ctrl+F.
2. **Shell Input Handling:**
  - Shell calls `send_signal_to_all(SIGFG)` ; to resume suspended processes.
3. **Inside `send_signal_to_all()`:**
  - Iterates over all processes.
  - For each process where `p->state == SUSPENDED`:
    - Sets `p->state = RUNNABLE`
    - Sets `p->suspended = 0`
4. **Result:** Suspended processes resume execution.

```

cpustat: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ test2
Parent started with pid: 3
Hello , I am pareChild started with pid: 4
Hi there , I am child
nt
Hello , I am parent
Hi there , I am child
Hello , I am parent
Hi there , I am child
Hello , I am parent
Hi there , I am child
Hello , I am parent
Hi there , I am child
Hello , I am parent
Hi there , I am child
Ctrl-B is detected by xv6
$ Ctrl-F is detected by xv6
Hello , I am parent
Hi there , I am child
Hello , I am parent
Hi there , I am child
Hello , I am parent

```

Figure 3: Ctrl+B and Ctrl + F

#### 4. Ctrl+G → SIGCUSTOM (Custom Signal)

1. **User Input:** User presses Ctrl+G.
2. **Shell Input Handling:**
  - Shell calls `send_signal_to_all(SIGCUSTOM)` ; to notify all processes.
3. **Inside `send_signal_to_all()`:**
  - Iterates over processes.
  - If a process has a custom handler (`p->signal_handler != 0`):
    - Sets `p->pending_signal = SIGCUSTOM`
    - If `p->state == SLEEPING`, wakes it up by setting `RUNNABLE`
4. **Result:** When scheduled, process checks for `pending_signal` and invokes its handler.

```

init: starting sh
$ test1
This is normal code running
This is normal code running
This is normal code running
This is normal code running
This is normal code running
This is normal code running
This is normal code running
This is normal code running
This is normal code running
This is normal code running
Ctrl-G is detected by xv6
I am inside the handler
I am Shivam
This is normal code running
This is normal code running
This is normal code running
This is normal code running
This is normal code running
This is normal code running

```

Figure 4: Ctrl+G

```

All child processes created with start_later flag set.
Calling sys_scheduler_start() to allow execution.
sys_scheduler_start called
Child 0 (PID: 5) started but should not run yet.
Child 1 (PID: 6) started but should not run yet.
Child 2 (PID: 7) started but should not run yet.
Child 0 (PID: 5) exiting.
PID: 5
TAT: 366
WT: 1
RT: 1
#CS: 255
Child 1 (PID: 6) exiting.
PID: 6
TAT: 677

```

Figure 5: Scheduler

## Dynamic Priority Scheduling: Effect of $\alpha$ and $\beta$ Parameters

The scheduling algorithm implemented uses a dynamic priority function to determine which process should be scheduled next. The priority  $\pi_i(t)$  of process  $P_i$  at time  $t$  is defined as:

$$\pi_i(t) = \pi_i(0) - \alpha \cdot C_i(t) + \beta \cdot W_i(t)$$

- $\pi_i(0)$ : Initial priority of process  $P_i$  (defined in `Makefile`)
- $C_i(t)$ : Total CPU time consumed by  $P_i$  up to time  $t$
- $W_i(t)$ : Waiting time of  $P_i$  since its creation
- $\alpha, \beta$ : Tunable parameters to control priority decay and boost (defined in `Makefile`)

### Interpretation of Parameters

- **Parameter  $\alpha$** : Controls how quickly a process loses priority based on its CPU usage.
  - A higher  $\alpha$  causes CPU-bound processes to lose priority faster.
  - Helps prevent long-running CPU-intensive tasks from hogging the CPU.
- **Parameter  $\beta$** : Controls how much waiting time contributes to the priority.
  - A higher  $\beta$  rewards processes that have waited longer.
  - Helps I/O-bound and short processes gain CPU access sooner.

## Effect on Different Process Types

Process Type	High $\alpha$ , Low $\beta$	Low $\alpha$ , High $\beta$
CPU-bound	Quickly penalized due to high CPU usage. Low chances of frequent scheduling.	Deprioritized slowly. May hog the CPU.
I/O-bound	May starve if $\beta$ is too low.	Frequently scheduled due to accumulated wait time.
Short jobs	May not get enough priority once CPU usage increases.	Scheduled quickly, finish early.

Table 1: Effect of  $\alpha$  and  $\beta$  on process behavior

## Experimental Observation

In the profiling experiments conducted:

- With higher  $\alpha$  (e.g., 3) and moderate  $\beta$  (e.g., 2):
  - CPU-bound jobs were penalized after initial execution.
  - Short/I/O-bound jobs were scheduled more frequently and completed quickly.
  - Fairness was observed with no starvation.
- With lower  $\alpha$  (e.g., 1) and higher  $\beta$  (e.g., 4):
  - Even minor waiting resulted in a large priority boost.
  - I/O-bound processes preempted CPU-bound ones very frequently.
  - Potential starvation of long CPU-intensive jobs was observed.
- With  $\alpha = 0$  and  $\beta = 0$ :
  - Priority remained static as  $\pi_i(t) = \pi_i(0)$ .
  - No dynamic adjustment based on CPU usage or waiting time.
  - Scheduling became unfair over time as processes with initially higher priority continued to dominate CPU access.
  - Both CPU-bound and I/O-bound processes risked starvation if they started with lower priority.

## Conclusion

The  $\alpha$  and  $\beta$  parameters significantly impact the system’s responsiveness and fairness:

- To favor interactive/I/O-bound jobs and prevent starvation, use a higher  $\beta$ .
- To prevent long-running CPU-bound tasks from dominating, use a higher  $\alpha$ .
- With  $\alpha = 0$  and  $\beta = 0$ , the scheduler becomes static, resulting in potential long-term unfairness.
- An optimal balance ensures all processes get CPU access proportionately based on their needs.
- To favor **interactive/I/O-bound jobs** and prevent starvation, use a **higher  $\beta$** .
- To prevent **long-running CPU-bound tasks from dominating**, use a **higher  $\alpha$** .
- An optimal balance ensures all processes get CPU access proportionately based on their needs.